

# Cumulative Prospect Theory Meets Reinforcement Learning: Prediction and Control

- **Author**
  - Prashanth L.A. et al.
- **Publication**
  - ICML 2016
- **Related Work**
  - Reinforcement Learning With Function Approximation for Traffic Signal Control
    - Prashanth L. A. & Shalabh Bhatnagar
    - IEEE T-ITS 2011
  - Threshold Tuning Using Stochastic Optimization for Graded Signal Control
    - Prashanth L. A. & Shalabh Bhatnagar
    - IEEE T-VT 2012
- [Cumulative Prospect Theory Meets Reinforcement Learning: Prediction and Control](#)
  - [Introduction](#)
  - [Traffic Light Control \(TLC\) Problem](#)
  - [CPT-value Estimation](#)
  - [Gradient-based Algo for CPT Optimization](#)
  - [Function Approximation and Boltzmann Policy](#)
  - [Experiments](#)

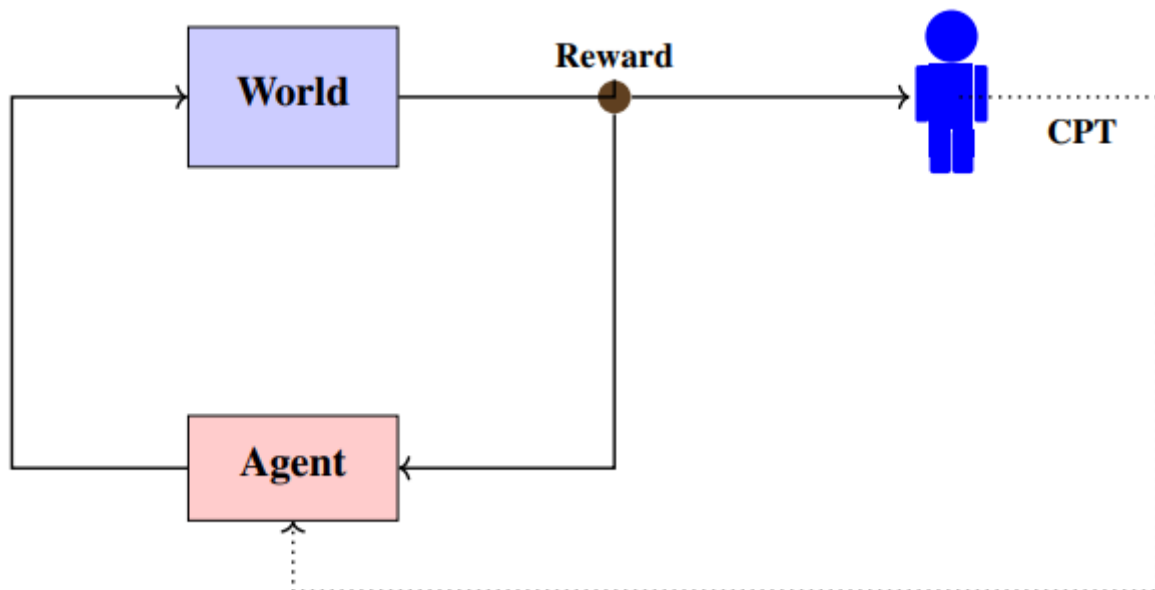
# Introduction

- **Contributions**

- The first to define and investigate *human-centered RL problem*.
- The first to combine CPT with RL.

- **Human-centered RL Problem**

- Agent controls a system to produce returns that are maximally aligned with the preferences of one or possibly multiple humans.
- Preferences of rational agents facing decisions with stochastic outcomes can be modeled using expected utilities.
- Here the authors use CPT to model the stochastic outcomes.



# Traffic Light Control (TLC) Problem

- **Target**
  - Train the traffic light to make a better traffic system
- **Environment**
  - Road network with signalled lanes that are spread across junctions and paths
  - Road users (cars)
- **Agent**
  - Traffic lights
- **States**
  - Queue length (pathwise)
  - Elapsed time (pathwise)
- **Actions**
  - feasible combinations of red and green
- **Reward**
  - CPT-value of differential delay  $X$  (because CPT needs gains and losses), i.e.,  $C(X) = \sum_{i=1}^M \mu_i C(X_i)$  where  $X_i$  is the differential delay (calculated by the elapsed time minus a baseline which is the elapsed time of a fixed-time signal control) of  $i$ -th path,  $\mu_i$  is the proportion of road users on the  $i$ -th path and  $M$  is the num of paths.

# CPT-value Estimation

- **Convergence assumptions (either or) for weighting function  $w$**

- Lipschitz continuous
- *Hölder* continuous

**Definition 1. (*Hölder continuity*)** A function  $f \in C([a, b])$  is said to satisfy a Hölder condition of order  $\alpha \in (0, 1]$  (or to be Hölder continuous of order  $\alpha$ ) if there exists  $H > 0$ , s.t.

$$\sup_{x \neq y} \frac{|f(x) - f(y)|}{|x - y|^\alpha} \leq H.$$

- Lipschitz continuity is the case that  $\alpha = 1$ .
- Locally Lipschitz

- **Estimation scheme under each of the assumptions**

- Under *Hölder* continuity, the CPT-value can be estimated by the **discrete** version.

$$\sum_{i=-m}^n \pi_i v(x_i),$$

where

$$\pi_i = \begin{cases} w(p_i + \dots + p_n) - w(p_{i+1} + \dots + p_n) & \text{for } 0 \leq i \leq n \\ w(p_{-m} + \dots + p_i) - w(p_{-m} + \dots + p_{i-1}) & -m \leq i < 0 \end{cases},$$

- **Algo:**

- 1: Simulate  $n$  i.i.d. samples from the distribution of  $X$ .
- 2: Order the samples and label them as follows:  $X_{[1]}, X_{[2]}, \dots, X_{[n]}$ . Note that  $u^+(X_{[1]}), \dots, u^+(X_{[n]})$  are also in ascending order.
- 3: Let

$$\overline{\mathbb{C}}_n^+ := \sum_{i=1}^n u^+(X_{[i]}) \left( w^+ \left( \frac{n+1-i}{n} \right) - w^+ \left( \frac{n-i}{n} \right) \right).$$

- 4: Apply  $u^-$  on the sequence  $\{X_{[1]}, X_{[2]}, \dots, X_{[n]}\}$ ; notice that  $u^-(X_{[i]})$  is in descending order since  $u^-$  is a decreasing function.
- 5: Let

$$\overline{\mathbb{C}}_n^- := \sum_{i=1}^n u^-(X_{[i]}) \left( w^- \left( \frac{i}{n} \right) - w^- \left( \frac{i-1}{n} \right) \right).$$

6: Return  $\overline{\mathbb{C}}_n = \overline{\mathbb{C}}_n^+ - \overline{\mathbb{C}}_n^-$ .

- Under Locally Lipschitz
  - Omitted

# Gradient-based Algo for CPT Optimization

- **Optimization Objective (Reward)**

$$\max_{\theta \in \Theta} C(X^\theta)$$

where the param vector  $\theta$  with dimension  $d$  is constrained in a compact and convex real set  $\Theta$ .

- **Gradient Estimation**

- It is hard to get the gradient of  $C(X^\theta)$ , especially in high-dimension cases.
- **SPSA** (Simultaneous Perturbation Stochastic Approximation)
  - A method to solve the problem above.
  - Use stochastic perturbation to get 2 sample values and approximate the gradient. The idea is similar to  $\frac{df(x)}{dx} \approx \frac{f(x+\varepsilon) - f(x-\varepsilon)}{2\varepsilon}$  where  $\varepsilon$  is a number close to 0.
  - At the  $n$ -th iteration, the gradient is estimated by

$$\hat{\nabla}_i C(X^\theta) = \frac{\bar{C}_n^{\theta_n + \delta_n \Delta_n} - \bar{C}_n^{\theta_n - \delta_n \Delta_n}}{2\delta_n \Delta_n^i}$$

where  $\bar{C}_n^{\theta_n}$  is the CPT-estimation calculated using the algo in *CPT-value Estimation* with  $m_n$  samples generated by param vector  $\theta_n$ ,  $\delta_n$  tends to 0 as  $n \rightarrow \infty$  (like  $\varepsilon$  above) and  $\Delta_n = (\Delta_n^1, \dots, \Delta_n^d)^\top$  where  $\{\Delta_n^i\}_{i=1}^d \stackrel{\text{i.i.d.}}{\sim} \text{Rademacher}$  (Rademacher distribution is 1 with half prob and  $-1$  with half prob).

- **Update Gradient**

- Update Rule (for the  $i$ -th component of  $\theta$ ):

$$\theta_{n+1}^i = \Gamma_i \left( \theta_n^i + \gamma_n \hat{\nabla}_i C(X^{\theta_n}) \right)$$

where  $\Gamma_i$  is the clip operator to constrain  $\theta$  in  $\Theta$  and  $\gamma_n$  is the learning rate (also called the step size).

- Convergence Condition
  - To guarantee that we can obtain the solution after several iterations, we need

$$\begin{cases} \gamma_n, \delta_n \rightarrow 0 \\ \frac{1}{m_n^{\frac{\alpha}{2}} \delta_n} \rightarrow 0 \\ \sum_n \gamma_n \rightarrow \infty \\ \sum_n \frac{\gamma_n^2}{\delta_n^2} < \infty \end{cases} \xRightarrow{\text{simple choice}} \begin{cases} \gamma_n = \frac{\gamma_0}{n} \\ m_n = m_0 n^\nu \text{ for } \begin{cases} \nu, \gamma > 0 \\ \gamma > \frac{\nu\alpha}{2} \end{cases} \\ \delta_n = \frac{\delta_0}{n^\gamma} \end{cases}$$

where  $\alpha$  is the *Hölder* order (we choose  $\alpha = 1$  for Lipschitz continuity).

- **Gradient Ascent Algo**

**Input:** initial parameter  $\theta_0 \in \Theta$  where  $\Theta$  is a compact and convex subset of  $\mathbb{R}^d$ , perturbation constants  $\delta_n > 0$ , sample sizes  $\{m_n\}$ , step-sizes  $\{\gamma_n\}$ , operator  $\Gamma : \mathbb{R}^d \rightarrow \Theta$ .

**for**  $n = 0, 1, 2, \dots$  **do**

Generate  $\{\Delta_n^i, i = 1, \dots, d\}$  using Rademacher distribution, independent of  $\{\Delta_m, m = 0, 1, \dots, n-1\}$ .

***CPT-value Estimation (Trajectory 1)***

Simulate  $m_n$  samples using  $(\theta_n + \delta_n \Delta_n)$ .

Obtain CPT-value estimate  $\overline{\mathbb{C}}_n^{\theta_n + \delta_n \Delta_n}$ .

***CPT-value Estimation (Trajectory 2)***

Simulate  $m_n$  samples using  $(\theta_n - \delta_n \Delta_n)$ .

Obtain CPT-value estimate  $\overline{\mathbb{C}}_n^{\theta_n - \delta_n \Delta_n}$ .

***Gradient Ascent***

Update  $\theta_n$  using the update rule above.

**end for**

**Return**  $\theta_n$ .

There still remains a problem: how the  $m_n$  samples are generated, i.e., what is the distribution (policy) we use? See the next chapter.



# Function Approximation and Boltzmann Policy

- **Function Approximation**

- Note that in  $Q$ -learning, the policy is dependent on the  $Q$ -table, which contains all  $Q(s, a)$  for each state  $s$  and each action  $a$ . At some state  $s$ , the agent choose the action with biggest  $Q$  value. The dimension of the  $Q$ -table is  $S \times A$  where  $S$  is the num of states and  $A$  is the num of feasible actions.
- In problems with high-dimension of states and actions,  $Q$ -table becomes computationally expensive. Thus, we use function approximation to approximate the  $Q$  funtion:

$$Q(s, a) \approx \theta^T \phi_{s, a}$$

where  $\phi_{s, a}$  is a  $d$ -dimensional vector describing the state-action feature, e.g.,  $d$  can be  $S + A$  if we construct features on each states and actions separately instead of considering their combinations.

- In this paper, the authors use the following settings of state-action feature  $\phi_{s, a}$ :

State	Action	Feature
$q_i(n) < L_1$ and $t_i(n) < T_1$	RED	0
	GREEN	1
$q_i(n) < L_1$ and $t_i(n) \geq T_1$	RED	0.2
	GREEN	0.8
$L_1 \leq q_i(n) < L_2$ and $t_i(n) < T_1$	RED	0.4
	GREEN	0.6
$L_1 \leq q_i(n) < L_2$ and $t_i(n) \geq T_1$	RED	0.6
	GREEN	0.4
$q_i(n) \geq L_2$ and $t_i(n) < T_1$	RED	0.8
	GREEN	0.2
$q_i(n) \geq L_2$ and $t_i(n) \geq T_1$	RED	1
	GREEN	0

where  $q$  represents the queue length,  $t$  represents the elapsed time and  $L_1, L_2, T_1$  are thresholds. Now the dimension  $d = 6A$  since they have shrinked the  $S$  states to be 6 states based on the thresholds.

- **Boltzmann Policy**

- Instead of choosing action with the biggest  $Q$  value, Boltzmann policy makes the  $Q$  values to be probs under some state.
- Formula:

$$\pi_{\theta}(s, a) = \frac{e^{\theta^{\top} \phi_{s, a}}}{\sum_{a' \in A(s)} e^{\theta^{\top} \phi_{s, a'}}}$$

where  $A(s)$  is the feasible action set under a certain state  $s$ .

- The exponential operator makes all values positive and thus can be transformed to be probs.


# Experiments

- **Experiment Design**


- CPT
  - Use CPT-value as reward.
- EUT
  - Use value function  $u$  but do not use weighting function  $w$ .
- AVG
  - Use neither value function nor weighting function, i.e., just use the simple mean.
- Do 100 independent tests, calculate CPT-value of these 3 algos respectively.

- **Results**


- CPT



- EUT



◦ AVG



- Results show that  $CPT > EUT > AVG$ . I think this is obvious and even need not be tested since the final measure is CPT-value and the CPT model is trained based on that.