

Cumulative Prospect Theory Meets Reinforcement Learning: Prediction and Control

- **Author**
 - Prashanth L.A. et al.
- **Publication**
 - ICML 2016
- **Related Work**
 - Reinforcement Learning With Function Approximation for Traffic Signal Control
 - Prashanth L. A. & Shalabh Bhatnagar
 - IEEE T-ITS 2011
 - Threshold Tuning Using Stochastic Optimization for Graded Signal Control
 - Prashanth L. A. & Shalabh Bhatnagar
 - IEEE T-VT 2012
- [Cumulative Prospect Theory Meets Reinforcement Learning: Prediction and Control](#)
 - [Introduction](#)
 - [Traffic Light Control \(TLC\) Problem](#)
 - [CPT-value Estimation](#)
 - [Gradient-based Algo for CPT Optimization](#)
 - [Function Approximation and Boltzmann Policy](#)
 - [Experiments](#)

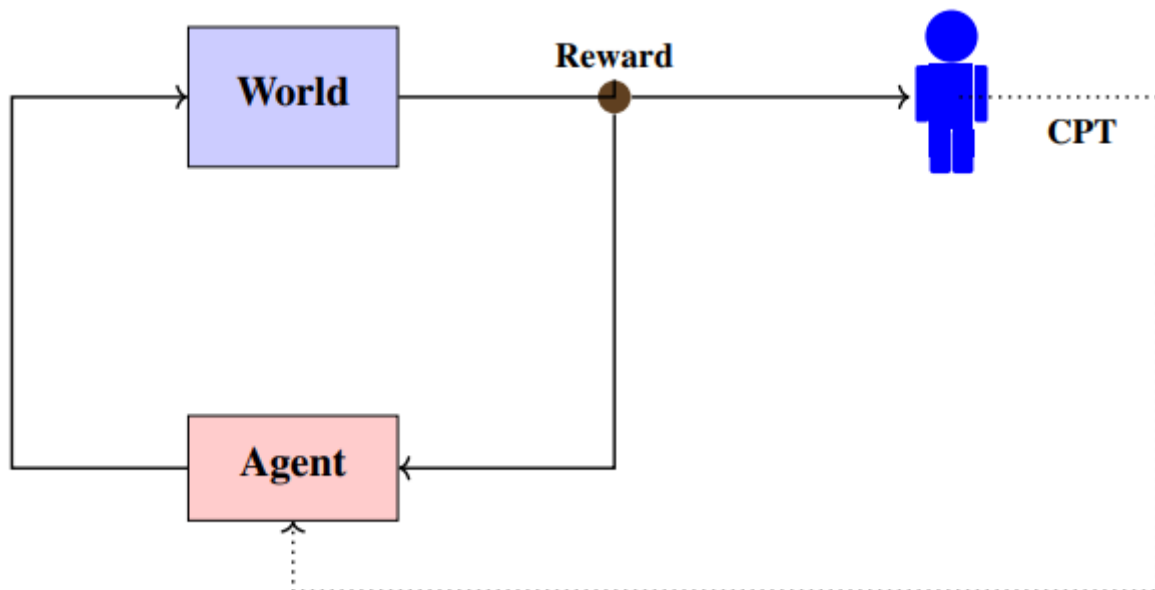
Introduction

- **Contributions**

- The first to define and investigate *human-centered RL problem*.
- The first to combine CPT with RL.

- **Human-centered RL Problem**

- Agent controls a system to produce returns that are maximally aligned with the preferences of one or possibly multiple humans.
- Preferences of rational agents facing decisions with stochastic outcomes can be modeled using expected utilities.
- Here the authors use CPT to model the stochastic outcomes.



Traffic Light Control (TLC) Problem

- **Target**
 - Train the traffic light to make a better traffic system
- **Environment**
 - Road network with signalled lanes that are spread across junctions and paths
 - Road users (cars)
- **Agent**
 - Traffic lights
- **States**
 - Queue length (pathwise)
 - Elapsed time (pathwise)
- **Actions**
 - feasible combinations of red and green
- **Reward**
 - CPT-value of differential delay X (because CPT needs gains and losses), i.e., $C(X) = \sum_{i=1}^M \mu_i C(X_i)$ where X_i is the differential delay (calculated by the elapsed time minus a baseline which is the elapsed time of a fixed-time signal control) of i -th path, μ_i is the proportion of road users on the i -th path and M is the num of paths.

CPT-value Estimation

- **Convergence assumptions (either or) for probability weighting function w**

- Lipschitz continuous
- *Hölder* continuous

Definition 1. (*Hölder continuity*) A function $f \in C([a, b])$ is said to satisfy a Hölder condition of order $\alpha \in (0, 1]$ (or to be Hölder continuous of order α) if there exists $H > 0$, s.t.

$$\sup_{x \neq y} \frac{|f(x) - f(y)|}{|x - y|^\alpha} \leq H.$$

- Lipschitz continuity is the case that $\alpha = 1$.
- Locally Lipschitz

- **Estimation scheme under each of the assumptions**

- Under *Hölder* continuity, the CPT-value can be estimated by the **discrete** version.

$$\sum_{i=-m}^n \pi_i v(x_i),$$

where

$$\pi_i = \begin{cases} w(p_i + \dots + p_n) - w(p_{i+1} + \dots + p_n) & \text{for } 0 \leq i \leq n \\ w(p_{-m} + \dots + p_i) - w(p_{-m} + \dots + p_{i-1}) & -m \leq i < 0 \end{cases},$$

- **Algo:**

- 1: Simulate n i.i.d. samples from the distribution of X .
- 2: Order the samples and label them as follows: $X_{[1]}, X_{[2]}, \dots, X_{[n]}$. Note that $u^+(X_{[1]}), \dots, u^+(X_{[n]})$ are also in ascending order.
- 3: Let

$$\overline{\mathbb{C}}_n^+ := \sum_{i=1}^n u^+(X_{[i]}) \left(w^+ \left(\frac{n+1-i}{n} \right) - w^+ \left(\frac{n-i}{n} \right) \right).$$

- 4: Apply u^- on the sequence $\{X_{[1]}, X_{[2]}, \dots, X_{[n]}\}$; notice that $u^-(X_{[i]})$ is in descending order since u^- is a decreasing function.
- 5: Let

$$\overline{\mathbb{C}}_n^- := \sum_{i=1}^n u^-(X_{[i]}) \left(w^- \left(\frac{i}{n} \right) - w^- \left(\frac{i-1}{n} \right) \right).$$

- 6: Return $\overline{\mathbb{C}}_n = \overline{\mathbb{C}}_n^+ - \overline{\mathbb{C}}_n^-$.

- Under Locally Lipschitz

- Omitted

Gradient-based Algo for CPT Optimization

- **Optimization Objective (Reward)**

$$\max_{\theta \in \Theta} C(X^\theta)$$

where the param vector θ with dimension d is constrained in a compact and convex real set Θ .

- **Gradient Estimation**

- It is hard to get the gradient of $C(X^\theta)$, especially in high-dimension cases.
- **SPSA** (Simultaneous Perturbation Stochastic Approximation)
 - A method to solve the problem above.
 - Use stochastic perturbation to get 2 sample values and approximate the gradient. The idea is similar to $\frac{df(x)}{dx} \approx \frac{f(x+\varepsilon) - f(x-\varepsilon)}{2\varepsilon}$ where ε is a number close to 0.
 - At the n -th iteration, the gradient is estimated by

$$\hat{\nabla}_i C(X^\theta) = \frac{\bar{C}_n^{\theta_n + \delta_n \Delta_n} - \bar{C}_n^{\theta_n - \delta_n \Delta_n}}{2\delta_n \Delta_n^i}$$

where $\bar{C}_n^{\theta_n}$ is the CPT-estimation calculated using the algo in *CPT-value Estimation* with m_n samples generated by param vector θ_n , δ_n tends to 0 as $n \rightarrow \infty$ (like ε above) and $\Delta_n = (\Delta_n^1, \dots, \Delta_n^d)^\top$ where $\{\Delta_n^i\}_{i=1}^d \stackrel{\text{i.i.d.}}{\sim} \text{Rademacher}$ (Rademacher distribution is 1 with half prob and -1 with half prob).

- **Update Gradient**

- Update Rule (for the i -th component of θ):

$$\theta_{n+1}^i = \Gamma_i \left(\theta_n^i + \gamma_n \hat{\nabla}_i C(X^{\theta_n}) \right)$$

where Γ_i is the clip operator to constrain θ in Θ and γ_n is the learning rate (also called the step size).

- Convergence Condition
 - To guarantee that we can obtain the solution after several iterations, we need

$$\begin{cases} \gamma_n, \delta_n \rightarrow 0 \\ \frac{1}{m_n^{\frac{\alpha}{2}} \delta_n} \rightarrow 0 \\ \sum_n \gamma_n \rightarrow \infty \\ \sum_n \frac{\gamma_n^2}{\delta_n^2} < \infty \end{cases} \xRightarrow{\text{simple choice}} \begin{cases} \gamma_n = \frac{\gamma_0}{n} \\ m_n = m_0 n^\nu \text{ for } \begin{cases} \nu, \gamma > 0 \\ \gamma > \frac{\nu\alpha}{2} \end{cases} \\ \delta_n = \frac{\delta_0}{n^\gamma} \end{cases}$$

where α is the *Hölder* order (we choose $\alpha = 1$ for Lipschitz continuity).

- **Gradient Ascent Algo**

Input: initial parameter $\theta_0 \in \Theta$ where Θ is a compact and convex subset of \mathbb{R}^d , perturbation constants $\delta_n > 0$, sample sizes $\{m_n\}$, step-sizes $\{\gamma_n\}$, operator $\Gamma : \mathbb{R}^d \rightarrow \Theta$.

for $n = 0, 1, 2, \dots$ **do**

Generate $\{\Delta_n^i, i = 1, \dots, d\}$ using Rademacher distribution, independent of $\{\Delta_m, m = 0, 1, \dots, n-1\}$.

CPT-value Estimation (Trajectory 1)

Simulate m_n samples using $(\theta_n + \delta_n \Delta_n)$.

Obtain CPT-value estimate $\overline{\mathbb{C}}_n^{\theta_n + \delta_n \Delta_n}$.

CPT-value Estimation (Trajectory 2)

Simulate m_n samples using $(\theta_n - \delta_n \Delta_n)$.

Obtain CPT-value estimate $\overline{\mathbb{C}}_n^{\theta_n - \delta_n \Delta_n}$.

Gradient Ascent

Update θ_n using the update rule above.

end for

Return θ_n .

There still remains a problem: how the m_n samples are generated, i.e., what is the distribution (policy) we use? See the next chapter.

Function Approximation and Boltzmann Policy

- **Function Approximation**

- Note that in Q -learning, the policy is dependent on the Q -table, which contains all $Q(s, a)$ for each state s and each action a . At some state s , the agent choose the action with biggest Q value. The dimension of the Q -table is $S \times A$ where S is the num of states and A is the num of feasible actions.
- In problems with high-dimension of states and actions, Q -table becomes computationally expensive. Thus, we use function approximation to approximate the Q funtion:

$$Q(s, a) \approx \theta^T \phi_{s, a}$$

where $\phi_{s, a}$ is a d -dimensional vector describing the state-action feature, e.g., d can be $S + A$ if we construct features on each states and actions separately instead of considering their combinations.

- In this paper, the authors use the following settings of state-action feature $\phi_{s, a}$:

| State | Action | Feature |
|---|--------|---------|
| $q_i(n) < L_1$ and $t_i(n) < T_1$ | RED | 0 |
| | GREEN | 1 |
| $q_i(n) < L_1$ and $t_i(n) \geq T_1$ | RED | 0.2 |
| | GREEN | 0.8 |
| $L_1 \leq q_i(n) < L_2$ and $t_i(n) < T_1$ | RED | 0.4 |
| | GREEN | 0.6 |
| $L_1 \leq q_i(n) < L_2$ and $t_i(n) \geq T_1$ | RED | 0.6 |
| | GREEN | 0.4 |
| $q_i(n) \geq L_2$ and $t_i(n) < T_1$ | RED | 0.8 |
| | GREEN | 0.2 |
| $q_i(n) \geq L_2$ and $t_i(n) \geq T_1$ | RED | 1 |
| | GREEN | 0 |

where q represents the queue length, t represents the elapsed time and L_1, L_2, T_1 are thresholds. Now the dimension $d = 6A$ since they have shrinked the S states to be 6 states based on the thresholds.

- **Boltzmann Policy**

- Instead of choosing action with the biggest Q value, Boltzmann policy makes the Q values to be probs under some state.
- Formula:

$$\pi_{\theta}(s, a) = \frac{e^{\theta^{\top} \phi_{s, a}}}{\sum_{a' \in A(s)} e^{\theta^{\top} \phi_{s, a'}}}$$

where $A(s)$ is the feasible action set under a certain state s .

- The exponential operator makes all values positive and thus can be transformed to be probs.


Experiments

- **Experiment Design**


- CPT
 - Use CPT-value as reward.
- EUT
 - Use value function u but do not use probability weighting function w .
- AVG
 - Use neither value function nor probability weighting function, i.e., just use the simple mean.
- Do 100 independent tests, calculate CPT-value of these 3 algos respectively.

- **Results**


- CPT



- EUT



○ AVG



- Results show that $CPT > EUT > AVG$. I think this is obvious and even need not be tested since the final measure is CPT-value and the CPT model is trained based on that.